



— AI Deep Search for  
Your Data

HybridRAG retrieval with multi-hop  
reasoning and citations

## PROBLEM

VRIN

20-30%

### Time Sink

Employees lose 20–30% of their week just finding information. McKinsey/IDC show **~1.8–2.5 hours/day** loss searching or re-creating info

[Read More ↗](#)

\$3.1 T/yr

### Economy Wide Losses

IBM/IDC report that the cost of “bad or inaccessible” data is massive. Gartner pegs poor data quality at **\$12.9–\$15 M/year per enterprise**

[Read More ↗](#)

20+

### SaaS tools & private stores

Teams use dozens of SaaS apps; answers require **cross-document, permission-aware, explainable reasoning** — today’s generic chatbots and **vanilla RAG fall short.**

## WHY NOW

VRIN

78%

### Mainstream Adoption

Companies report using AI in at least one function (up from 72% in 2024)

[Read More ↗](#)

92%

### Budget Rise

Executives expect to increase AI spend over the next 3 years; pressure to prove ROI is high

[Read More ↗](#)

60%

### SMB Momentum

Will increase spending in 2025 (faster than many other IT categories)

[Read More ↗](#)

100%

### Market Signal

Market leaders center messaging around permission-aware search, grounded answers — demand for cited, context-rich systems is durable

From documents to decisions.

Reasoned answers, permission-aware, grounded in company data

## 1. Facts-First Approach

 Extracted 15 facts and 7 vector chunks

- Knowledge Graphs + Vector Chunks to ground every response in company facts
- NO Hallucinations

## 2. Multi-hop Reasoning

 Explored 3 Source Documents

- Cross-document synthesis
- Traces the chain of facts
- Produces citations and confidence

## 3. BYOC/BYOK

- Respect ACLs
- Deploy in customer cloud
- Complete audit trail

## 4. Expert Specialization

- Custom experts for Support, Engineering, Legal, Finance & other industries

## MARKET OPPORTUNITY

A wedge into a  
\$7B core —  
expanding into  
the broader  
GenAI for Work  
stack

[Read More ↗](#)

TAM

Global enterprise/workplace search

\$6.97B

SAM

78% orgs willing to adopt  
AI in search now

\$5.44B

SOM

20% of SAM

\$1.1 B

```
from vrin import VRINClient

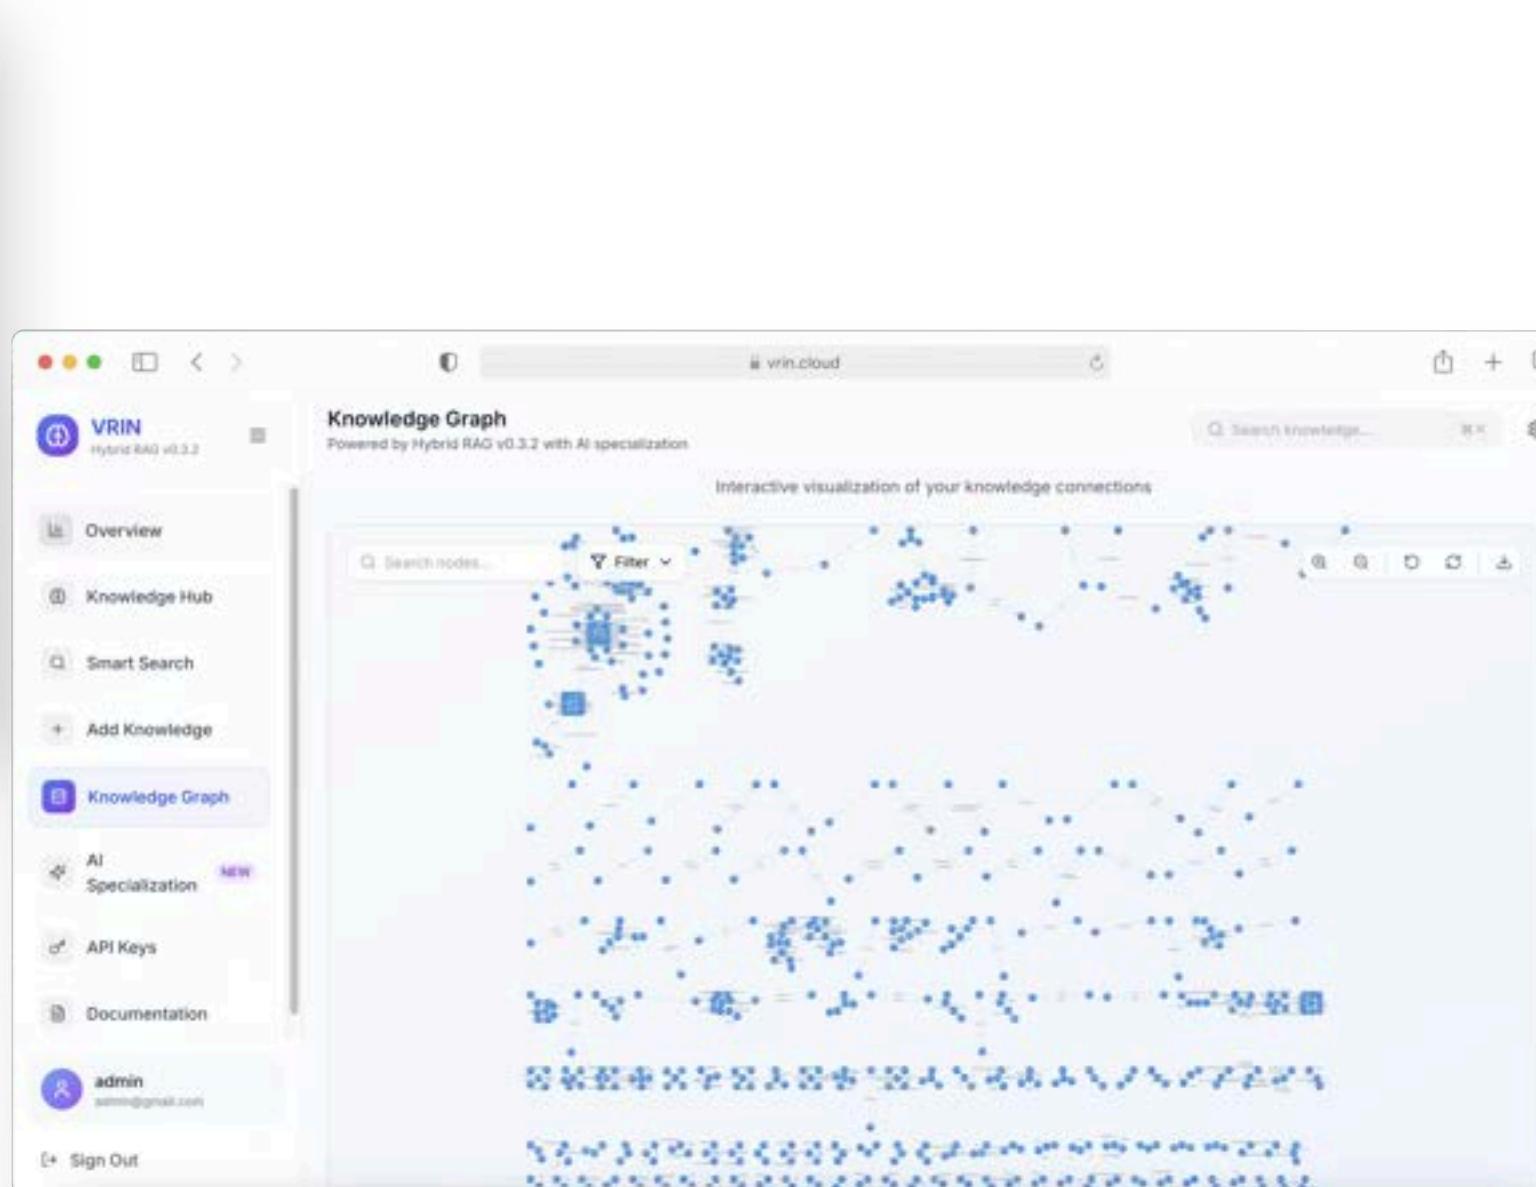
client = VRINClient(api_key=os.getenv("API_KEY"))

insert_doc = client.insert(
    content=tech_acquisition_doc,
    title="TechGlobal Strategic M&A Assessment"
)
```

## Ingest

---

Connectors → de-dup →  
extract entities/relations  
and embeddings



## Structure

---

Store as knowledge  
graph + vector index.

```
query = """
Analyze the TechGlobal-EuroManufacturing acquisition. Find
me conclusive evidence that justifies the $12B price tag,
and how does this
acquisition give TechGlobal competitive advantages over
QuantumCorp?
"""

response = client.query(query)
summary = response['summary']

# Other performance metrics
total_facts = response['total_facts']
total_chunks = response['total_chunks']
entities = response['entities']
multi_hop_chains = response['multi_hop_chains']
cross_doc_patterns = response['cross_document_patterns']
reasoning_confidence = response['reasoning_confidence']
```

## Retrieve & Reason

---

Entity + semantic retrieval → result  
fusion → multi-hop chains →  
answers with citation

# PRODUCT DEMO

VRiN

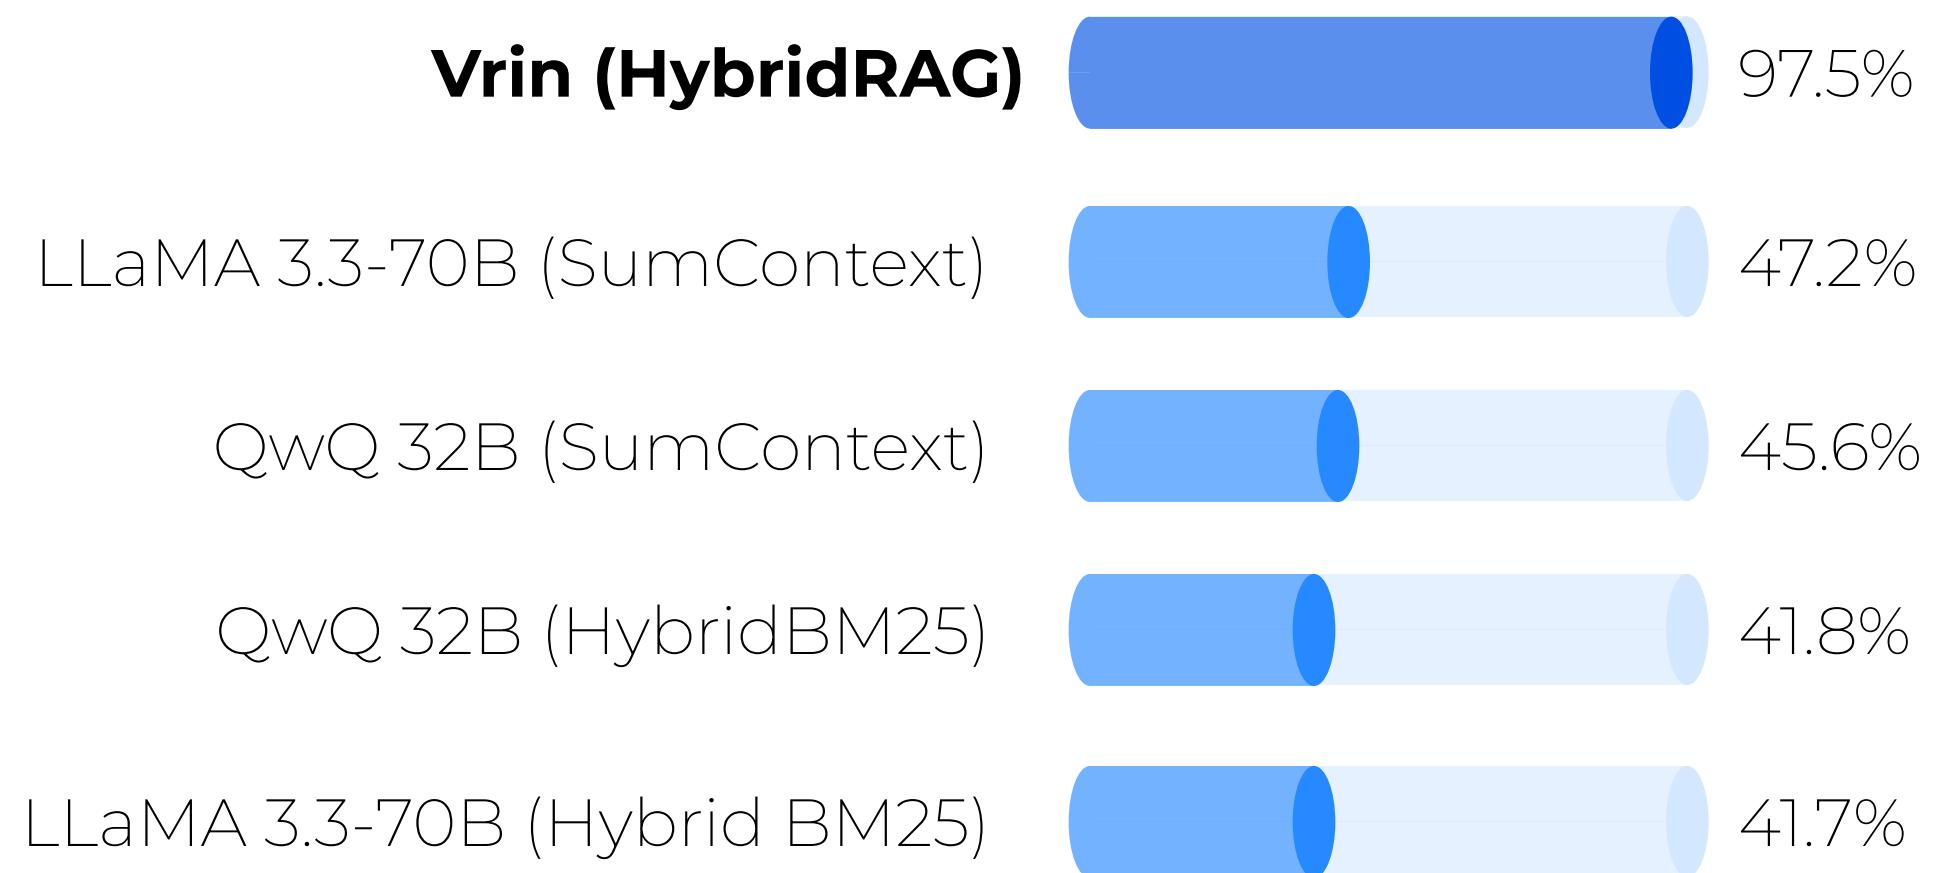
The screenshot shows a web-based application interface for VRiN. The top navigation bar includes tabs for "Overview", "Devices", "Knowledge Graph", "Event Stream", "AI Operations", "API Keys", "Dashboards", and "Logs". The main content area is titled "Overview" and features a prominent blue banner with the message "Welcome back, admin! 🎉 Your knowledge graph is growing. Here's what's happening today." Below the banner, there are four key metrics displayed in colored boxes: "Knowledge Nodes" (0), "Connections" (0), "Events (Processed)" (0), and "Data Centers" (0.000). A "Quick Actions" section contains four buttons: "Add Knowledge" (Import new information and start defining), "Scan Device" (Every 10 minutes via Device Log), "Explore deeper" (Discover more knowledge from logs), and "Go Expert" (Configure context AI preferences).

Watch ↗

## BENCHMARKS

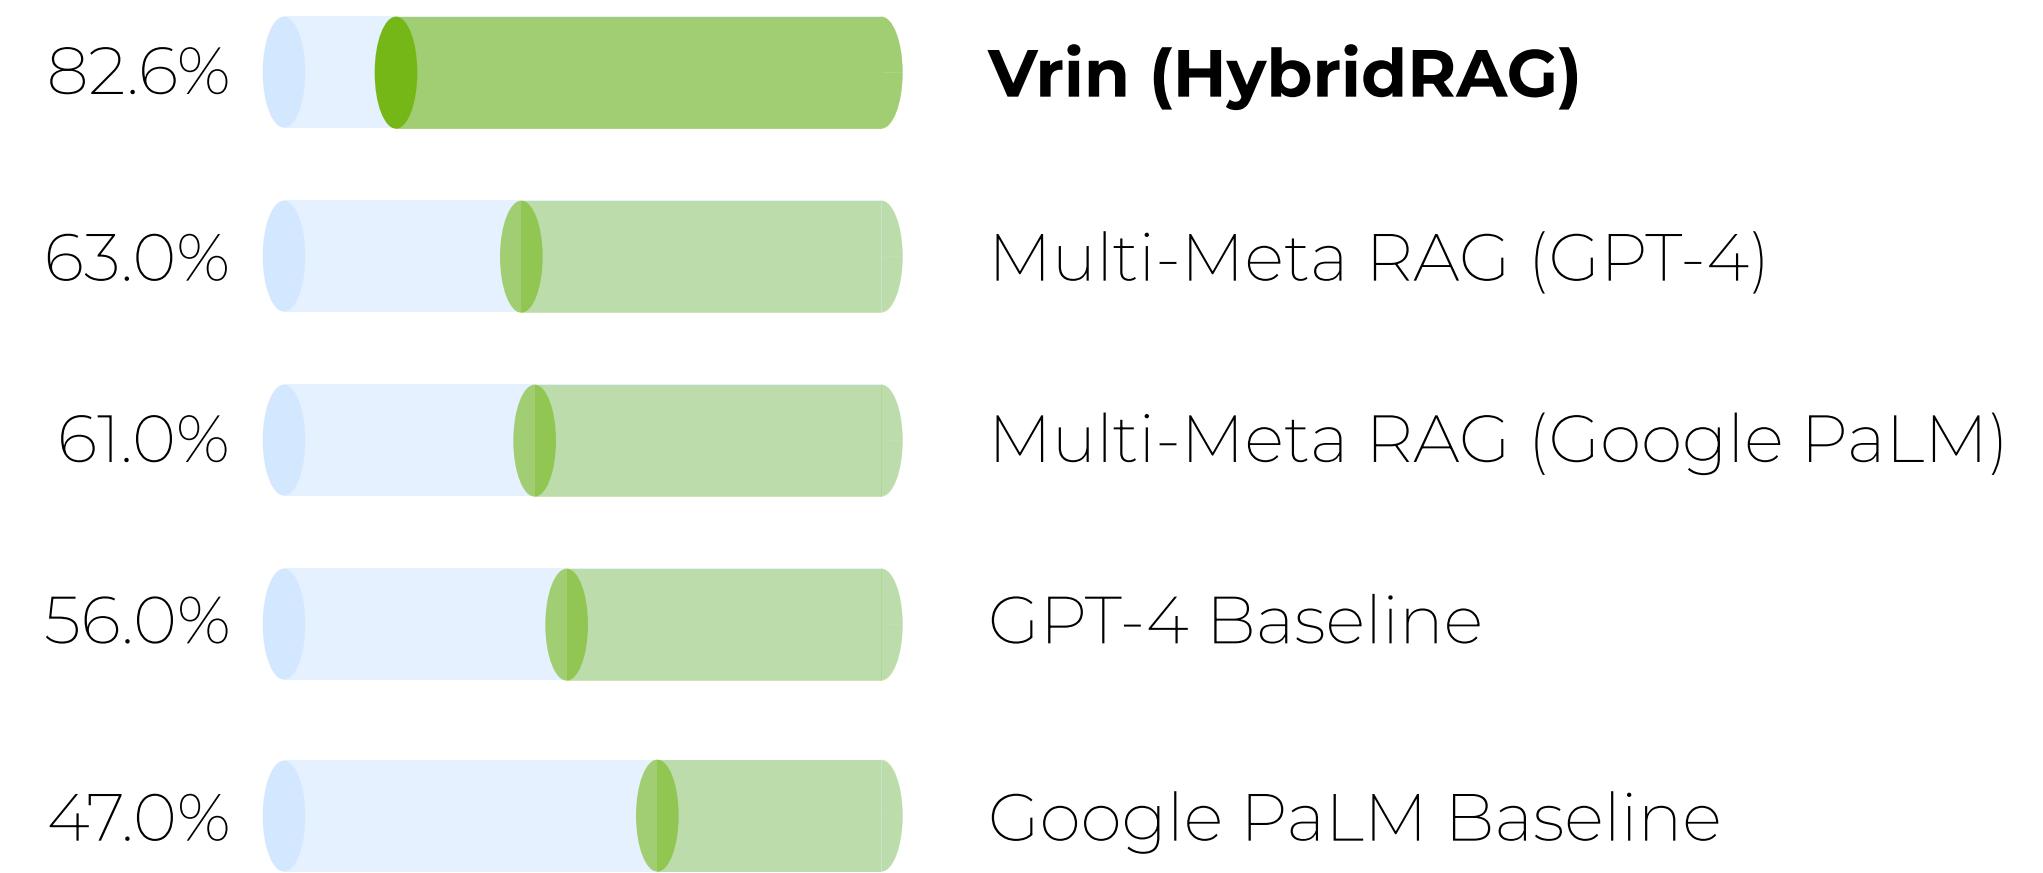
## RAGBench FinQA

Metric: Number Match (NM)



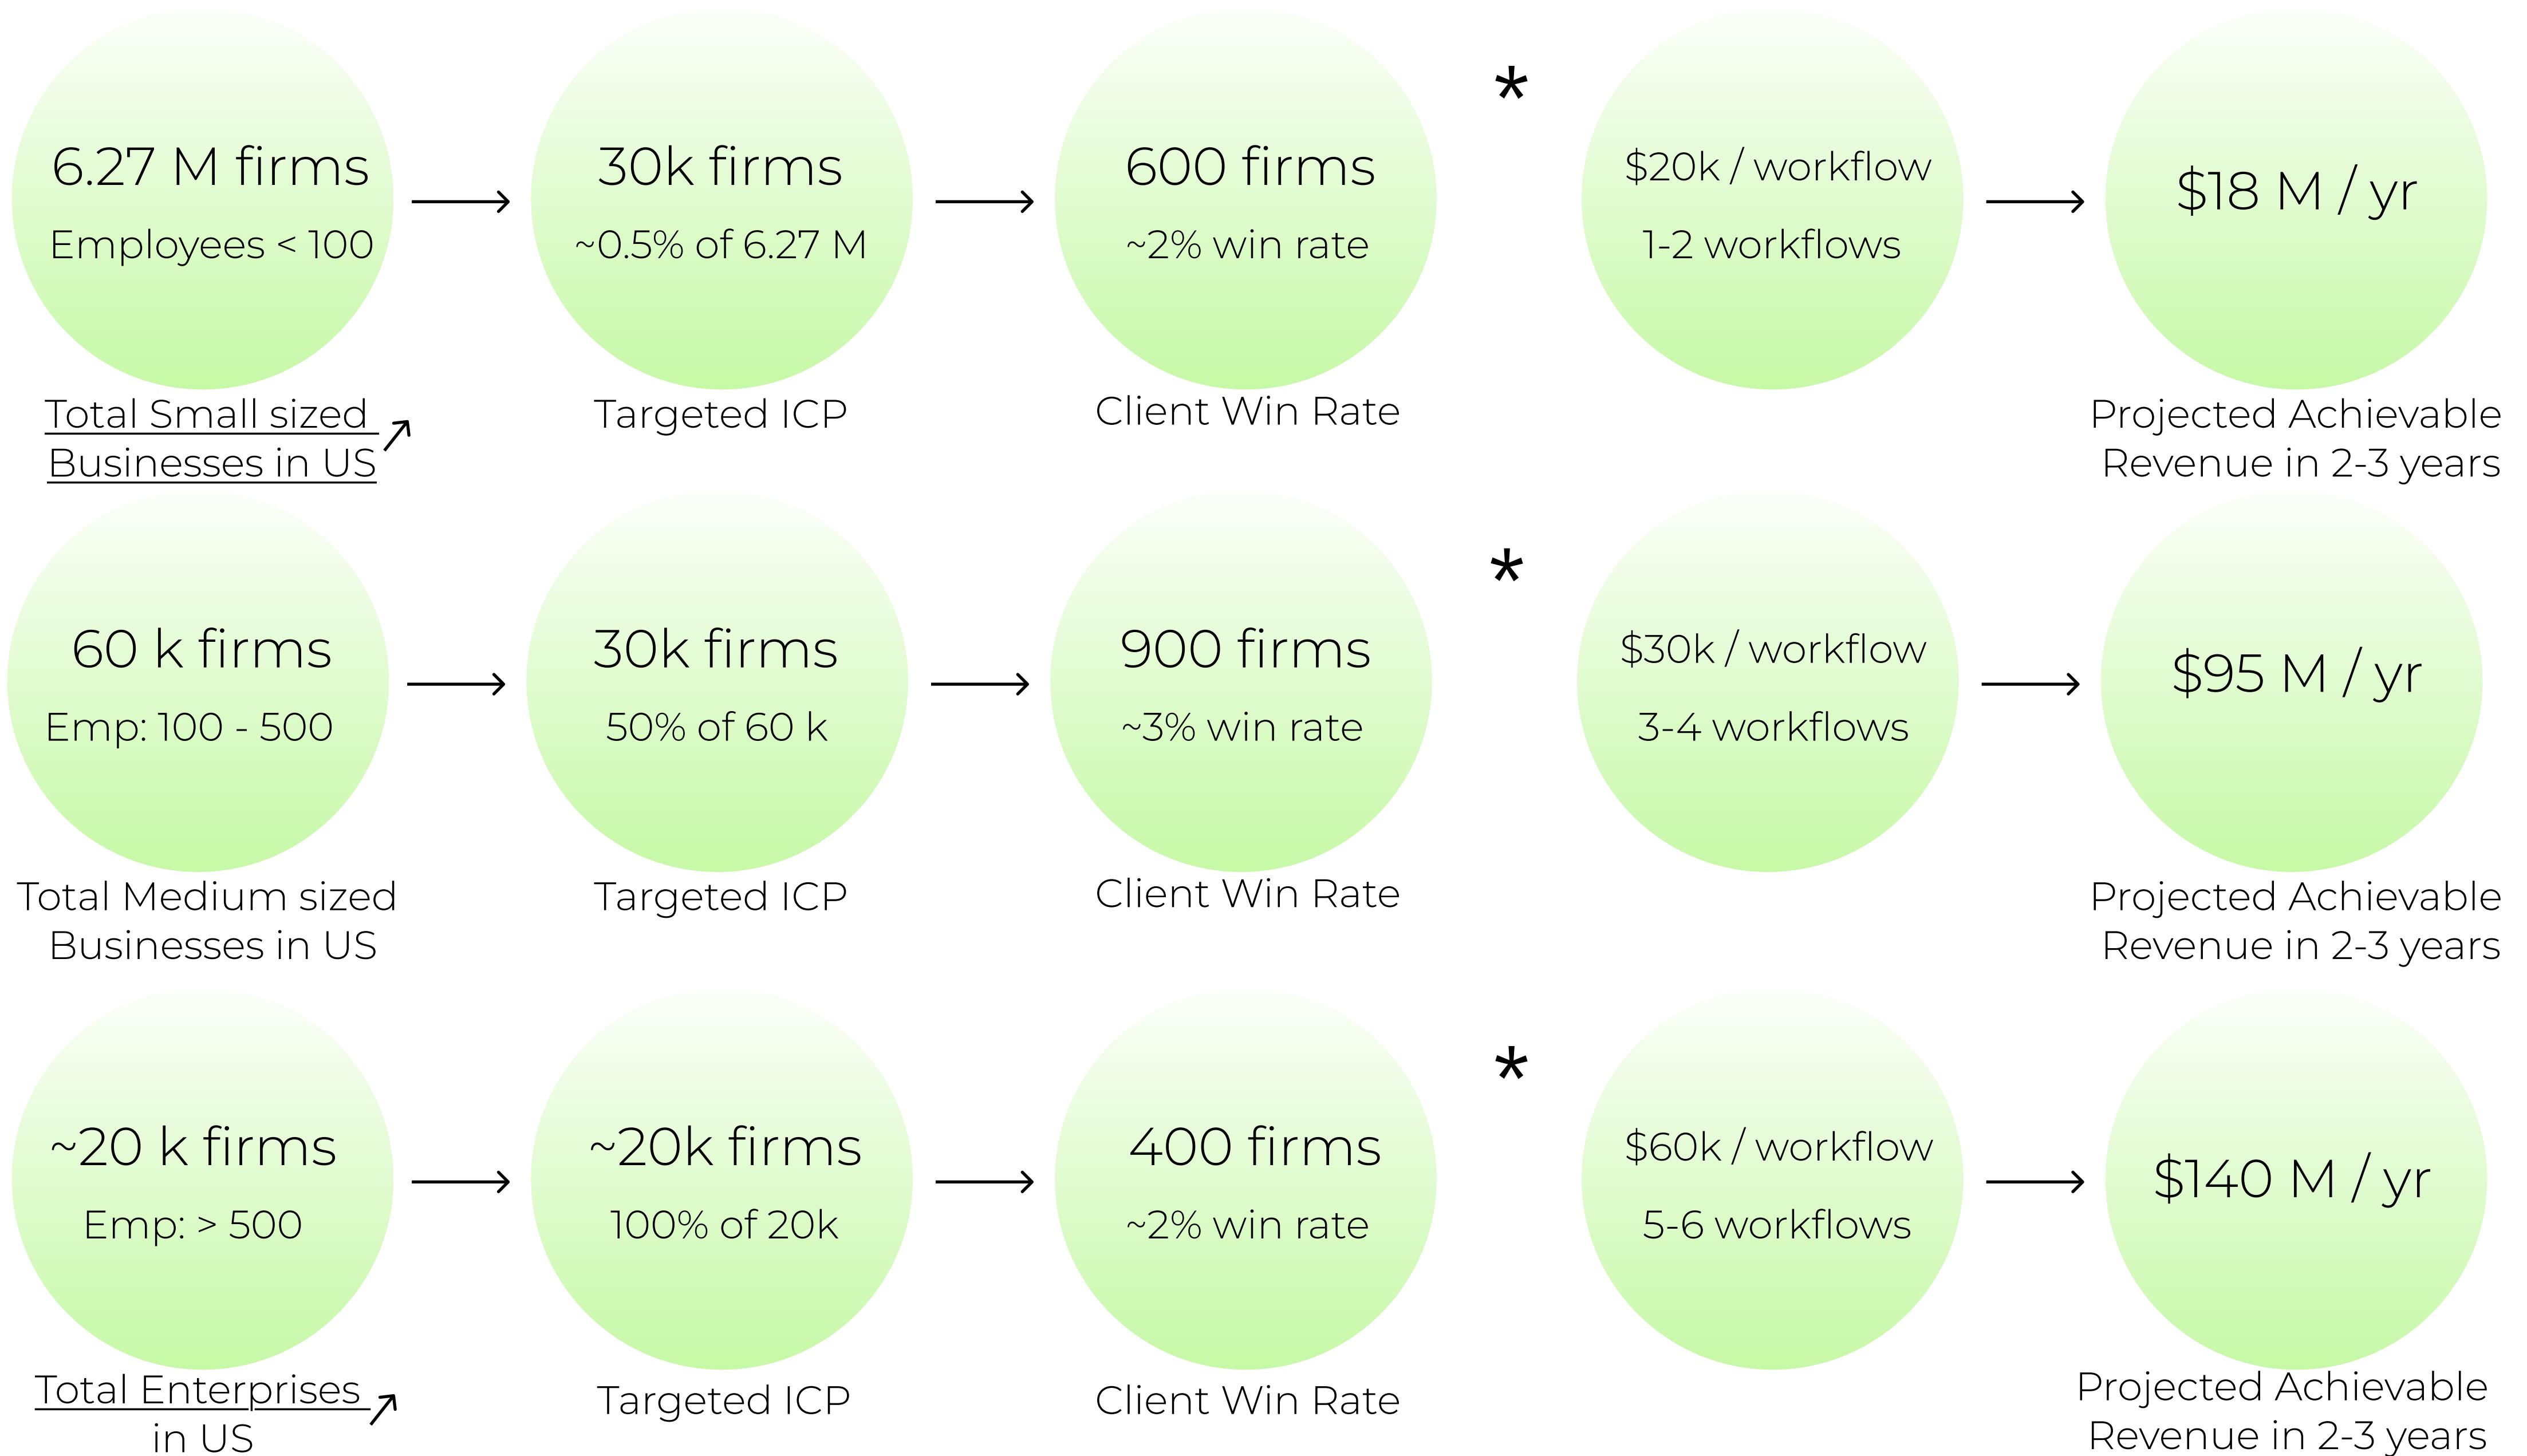
## MultiHop-RAG

Metric: Semantic Accuracy (SA)

[Learn More ↗](#)[Learn More ↗](#)

# REVENUE MODEL

VRiN



# BUSINESS MODEL

VRiN

Features	Builder (Free)	Team (\$2–4k/mo)	Business (\$8–15k/mo)	Enterprise (\$60–\$250k+/yr + usage)
Data scope	up to 200k chunks / 200k edges	2–5M chunks / 3–7.5M edges	10–25M chunks / 15–35M edges	Custom (100M+ chunks; 150M+ edges)
Queries/mo	10k	100–250k	500k–1M	Custom (SLA'd)
HybridRAG	Shared infra	Dedicated indices	Dedicated + VPC peering	Private/VPC or on-prem
Memory & CBOM	Basic	Full CBOM & TTL	Full + exports	Full + auditor packs
Security	API keys	Basic RBAC	SSO/SAML + SCIM	SSO/SAML, SCIM, data residency
Connectors	CSV/S3	+ Postgres/Drive	+ Slack/Jira/Confluence/SharePoint	All + custom
Support	Community	Email (48h SLA)	Priority (8–12h SLA)	TAM & DSE
Add-ons	—	Extra storage/queries	Compliance exports, private LLM	On-prem, managed upgrades

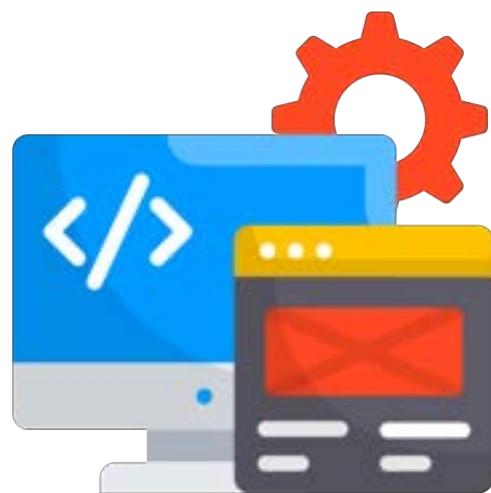
# GO-TO-MARKET

VRIN



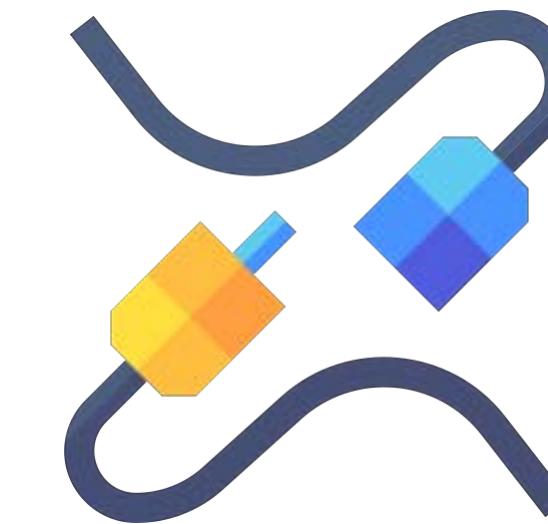
## Founder-Led Enterprise Sales

Target **CIO/CDO/Head of Platform** at data-rich, compliance-bound orgs (engineering, financial services, healthcare, industrial). Lead with pilot that measures time-to-answer and reduction in duplicate work.



## Self Serve Developer Motion

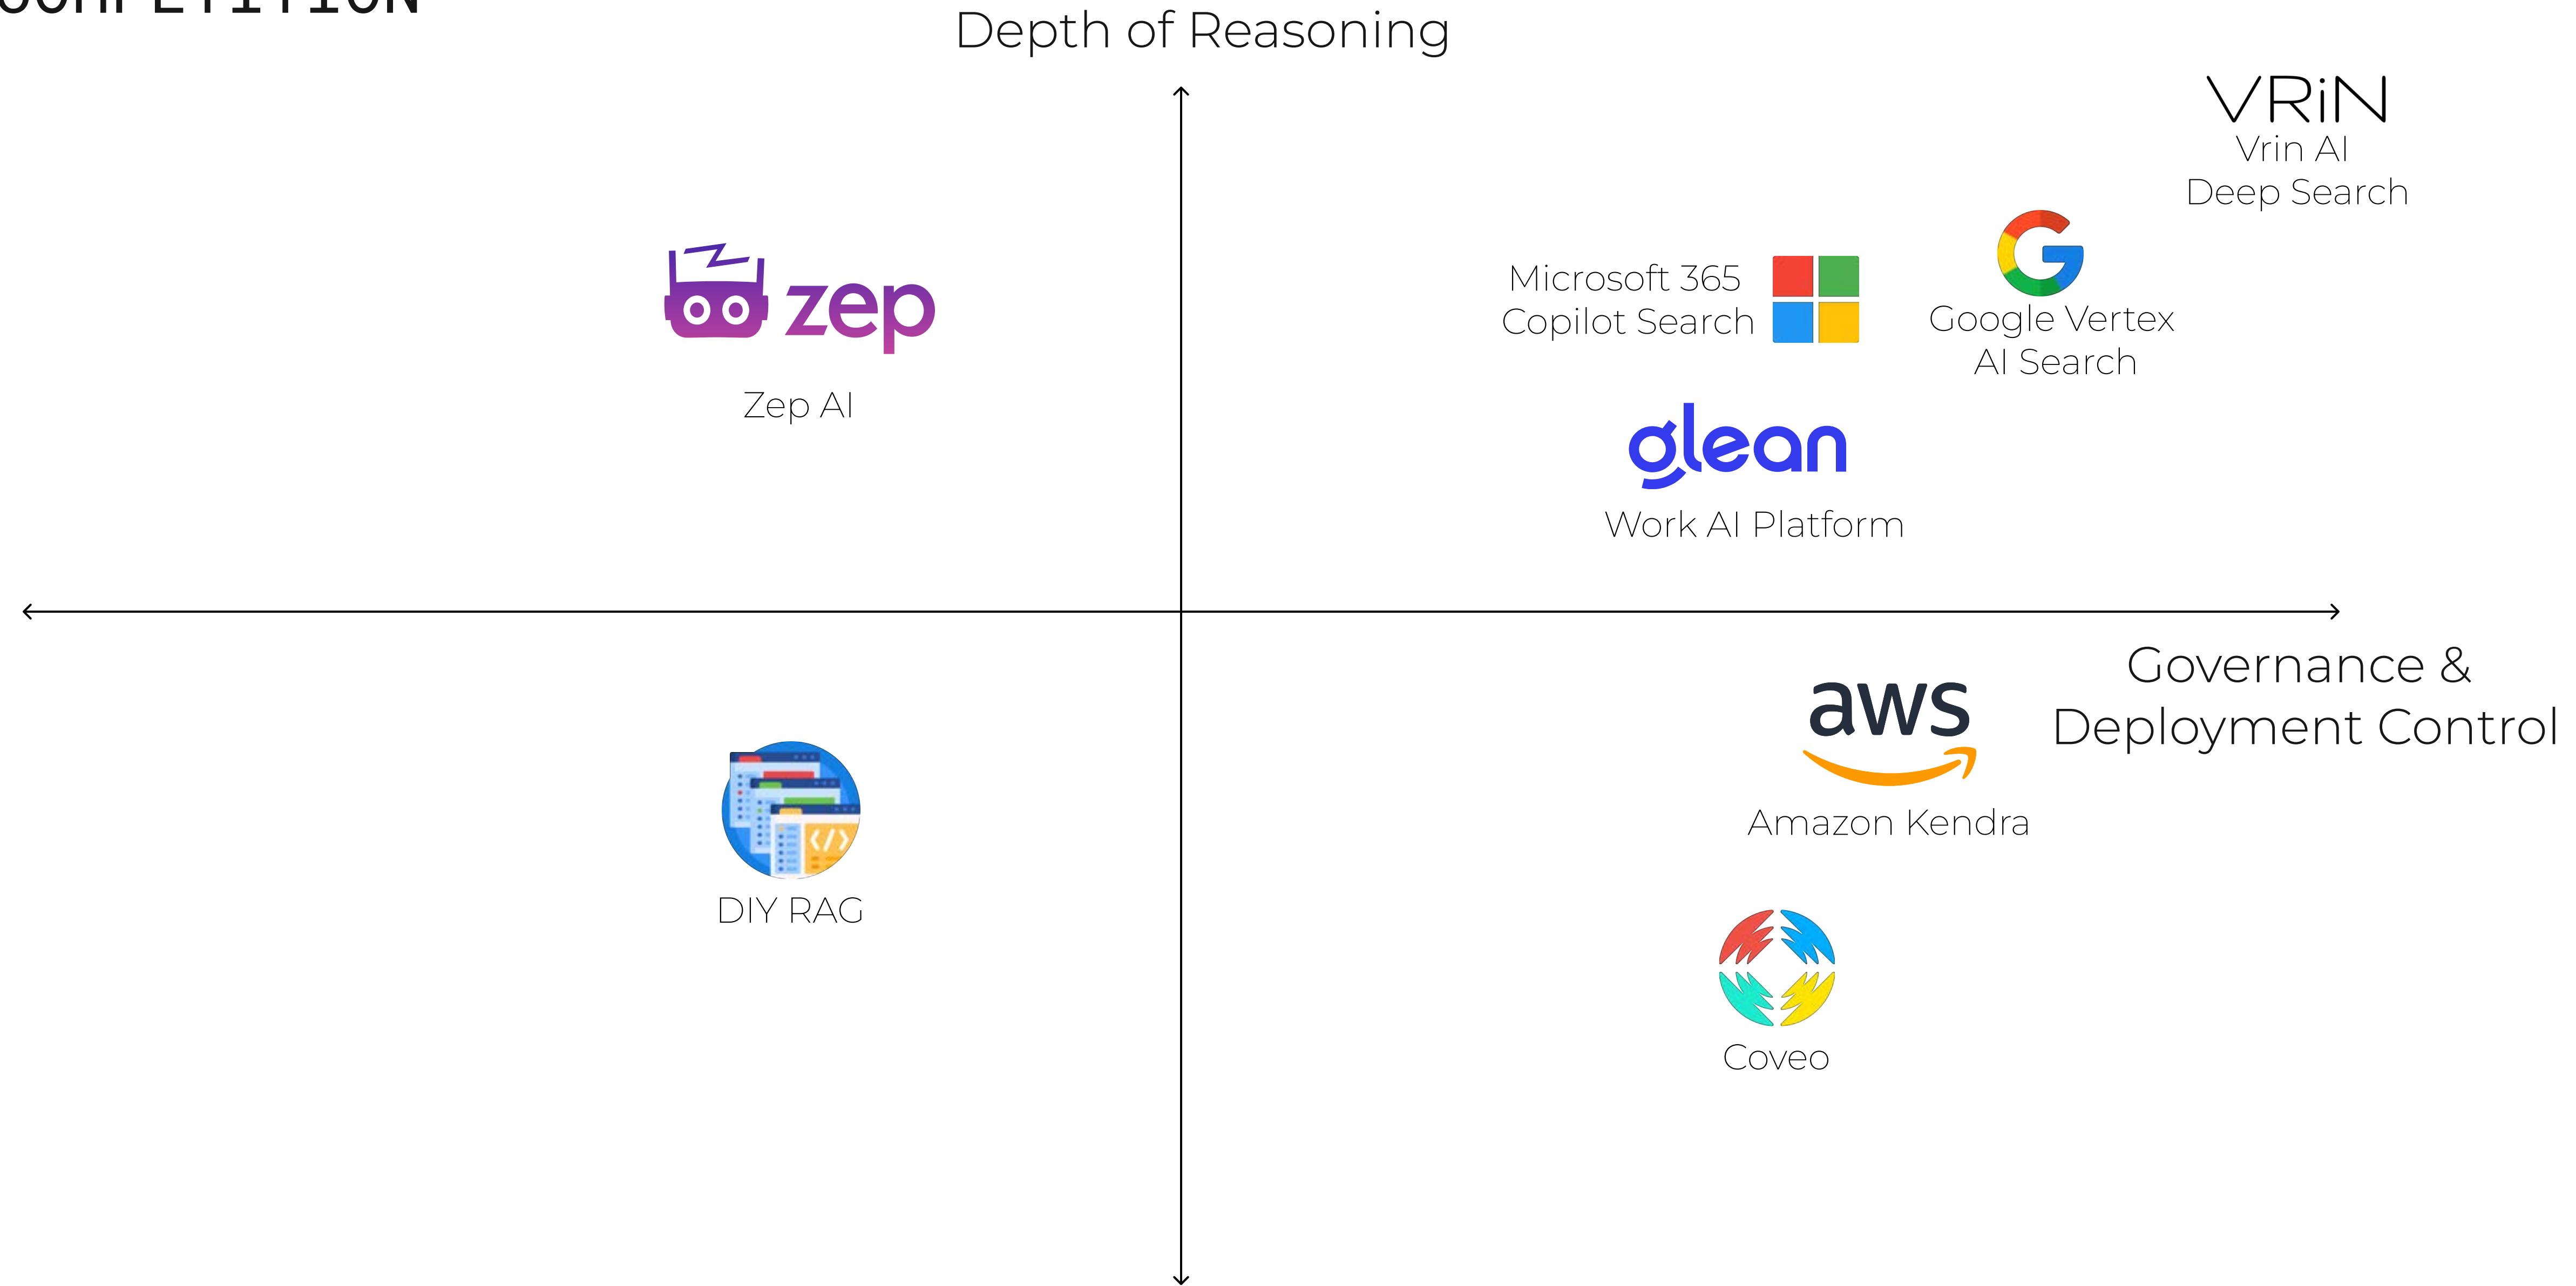
SDKs  
Guides  
Learning Modules  
Tutorials  
Community



## SaaS tools & private stores

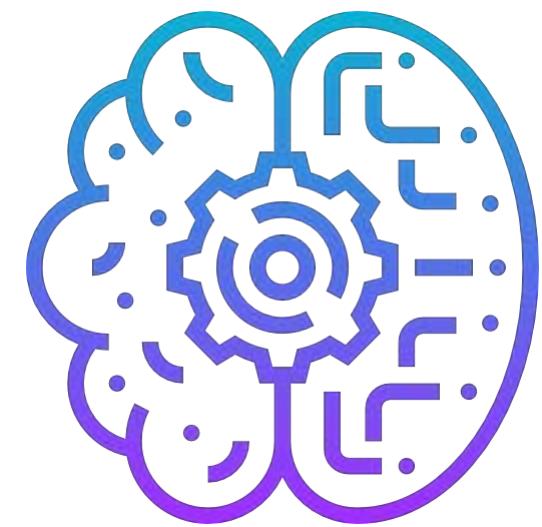
iPaaS & data integration partners for connectors. Connectors are commodities, reasoning & analysis are the real value

# COMPETITION



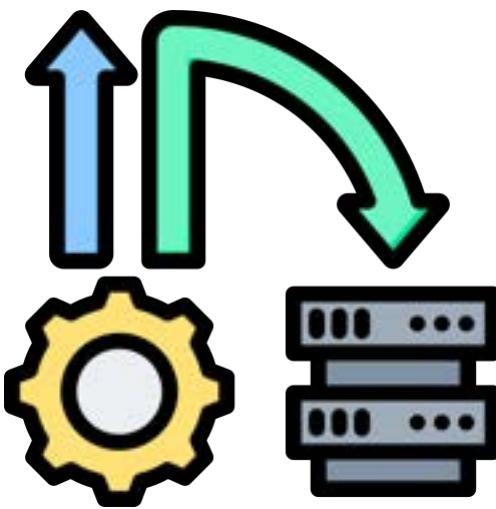
# COMPETITIVE ADVANTAGE

VRIN



## Deeply reasoned Explainable responses

Knowledge Graphs + Vectors to support multi-hop reasoning while retaining permission awareness to address industry's need for grounded outputs and lower hallucinations.



## Governance & deployment control as first-class

Similar governance posture as hyper-scalers (AWS Kendra, Google Vertex) but deeper cross-source reasoning chains and domain specializations instead of generic retrieval.



## Faster Path to ROI

We optimize for answer quality with citations and time-to-value.  
“Trust + Measurable Impact” wins budget in 2025.  
Ship deep search/answers inside existing tools via clean APIs & SDKs for rapid integration



## TEAM

I built a **production AI assistant system** for internal developers at **Microsoft** and saw vanilla RAG break under real workflows—fragile pipelines, context bloat, and no provenance. At UC Davis, my research pushed TinyML ECG models to extend battery life from ~2 weeks to 1+ month and used cloud offload for heavier multi-label analysis (**research published in IEEE ISCAS**). I also led a multi-agent Health LLM system that had to fetch the right medical signals and choose the right model for each question—exposing **RAG's gaps in multi-hop reasoning, cross-document synthesis, and time awareness**. Vrin is the system I wished I'd had: a neutral control plane that makes AI answers accurate, auditable, and enterprise-ready.

[Learn More ↗](#)

# TRACTION



## The Challenge

A leading research lab at UC Davis is standing up a multi-agent Health LLM to answer patient questions using continuously recorded signals plus clinical context. The existing approach struggled with reasoning — vector-only RAG missed cross-signal links and couldn't explain answers; and integration velocity—DIY pipelines across silos were slow to deploy and hard to maintain.

## The Solution

We integrated Vrin as the memory & context layer inside the multi-agent orchestration, providing HybridRAG with permission-aware retrieval, per-claim citations, and BYOK/VPC deployment patterns aligned to the health system's AI/privacy guidance.



## The Challenge

A multi-brand hospitality & retail operator needed a smarter website assistant. Knowledge was spread across a proprietary customer-insights database plus wikis, PDFs, and ticket notes — each with strict permissions. The team wanted trustable responses that referenced their own data, without standing up a costly DIY RAG stack or rewriting existing systems.

## The Solution

We deployed Vrin via SDK into their existing web assistant. Customer-insights tables and content were ingested with entity/relationship extraction, then fused during retrieval to support multi-hop, cited answers. Row-level permissions and VPC/BYOK patterns enforced data access. The assistant calls Vrin's API for grounded responses.



**Pipeline: Engaged with & Actively scoping multiple enterprise and mid-market buyers; pilot statements of work in progress to embed Vrin in priority workflows.**

## THE ASK

We're raising 12 months financing for GTM, Security/Compliance, iPaaS/connector partners, Reasoning & Governance R&D to land pilots → and convert to scalable motion.



\$500K

# VRin

Thank you



Vedant Patel



vedant@vrin.cloud



+1 (530)-204-8133

<https://www.vrin.cloud>

The screenshot shows a web browser window displaying the VRin cloud homepage. The URL 'vrin.cloud' is visible in the address bar. The page features a large, light blue background with a subtle grid pattern. In the center, the text 'AI Deep Search for your Data' is displayed, with 'your Data' in a different color (blue). Below this, the text 'HybridRAG retrieval with multi-hop reasoning and citations' is shown. At the bottom, there are two buttons: 'Get Started' and 'Schedule a Demo'. The top navigation bar includes links for 'Demo', 'Benchmarks', 'Industries', and 'Enterprise'. A small 'Now Available' button is located in the upper right corner of the main content area.

Trusted by Developers at

UCDAVIS  
UNIVERSITY OF CALIFORNIA